

3. Core Relational Concepts (Tables, Keys, Normalization)

Think of this section as your “SQL DNA class.” We’re going to cover what makes data relational, why primary and foreign keys matter, and how normalization shapes your data — all the stuff that makes your queries reliable and your pipelines sane.

3.1 What is a Table?

In SQL, a **table** is a **relation**, which is a fancy math word for a set of tuples (rows). Important things to remember:

1. **Rows are unordered** — SQL only guarantees order if you use `ORDER BY`.
2. **Columns have types** — the schema defines data type, nullability, default values, etc.
3. **Tables are sets of rows** — duplicates are possible unless restricted by constraints.

Example Table — Customers

```
CREATE TABLE dbo.Customers (
    CustomerID INT IDENTITY PRIMARY KEY,
    FirstName NVARCHAR(50) NOT NULL,
    LastName NVARCHAR(50) NOT NULL,
    Email NVARCHAR(100) UNIQUE,
    CreatedAt DATETIME2 DEFAULT SYSUTCDATETIME()
);
```

- `CustomerID` → Primary key (unique row identifier)
 - `Email` → Unique constraint (no duplicate emails)
 - `CreatedAt` → Default value
 - `NVARCHAR` → Unicode strings
-

3.2 Primary Key (PK)

A **primary key** uniquely identifies a row in a table.

Rules & Tips:

- Cannot be `NULL`.
- Can be single column or composite (multiple columns).
- Often used in indexes to speed up joins.

Example — Single PK

```
CREATE TABLE dbo.Products (
    ProductID INT IDENTITY PRIMARY KEY,
    ProductName NVARCHAR(100) NOT NULL,
    Price DECIMAL(10,2) NOT NULL
);
```

Example — Composite PK

```
CREATE TABLE dbo.OrderItems (
    OrderID INT NOT NULL,
    ProductID INT NOT NULL,
    Quantity INT NOT NULL,
    UnitPrice DECIMAL(10,2) NOT NULL,
    PRIMARY KEY (OrderID, ProductID)
);
```

- Composite PK = combination of columns must be unique.
- Useful in **junction tables** (many-to-many relationships).

3.3 Foreign Key (FK)

A foreign key enforces referential integrity. It ensures a column or set of columns **references a valid primary key** in another table.

Example — Orders referencing Customers

```
CREATE TABLE dbo.Orders (
    OrderID INT IDENTITY PRIMARY KEY,
    CustomerID INT NOT NULL,
    OrderDate DATETIME2 NOT NULL DEFAULT SYSUTCDATETIME(),
    CONSTRAINT FK_Orders_Customers FOREIGN KEY (CustomerID)
        REFERENCES dbo.Customers(CustomerID)
);
```

- `CustomerID` in `Orders` **must exist** in `Customers`.
- Helps prevent “orphan” rows.

Cascading Options

- `ON DELETE CASCADE` → delete child rows if parent is deleted
- `ON UPDATE CASCADE` → update child if parent key changes

```
ALTER TABLE dbo.Orders
ADD CONSTRAINT FK_Orders_Customers
FOREIGN KEY (CustomerID)
REFERENCES dbo.Customers(CustomerID)
ON DELETE CASCADE;
```

3.4 Data Relationships (ER perspective)

- One-to-Many (1:N) — One customer → many orders.
- Many-to-Many (M:N) — Many products → many orders → junction table (`OrderItems`).
- One-to-One (1:1) — Rare; e.g., User → UserProfile.

Visual example:

```
Customers (1) —< Orders (N) —< OrderItems (N) >- Products (1)
```

3.5 Normalization — Why we care

Normalization is the art of reducing redundancy and anomalies.

1NF (First Normal Form)

- Each column is atomic (no lists or arrays).
- Each row is unique.

```
-- ✗ Bad (violates 1NF)
CREATE TABLE dbo.BadCustomerOrders (
    CustomerID INT,
    Orders NVARCHAR(MAX) -- list of order IDs
);
```

```
-- ☑ Good
CREATE TABLE dbo.Orders (
    OrderID INT IDENTITY PRIMARY KEY,
    CustomerID INT NOT NULL
);
```

2NF (Second Normal Form)

- Table is in 1NF.

- Every non-PK column depends on the whole PK (applies to composite keys).

```
-- OrderItems table example
PRIMARY KEY (OrderID, ProductID)
-- Quantity depends on both OrderID and ProductID, not just one.
```

3NF (Third Normal Form)

- Table is in 2NF.
- No transitive dependencies: non-PK columns do not depend on other non-PK columns.

```
-- ✗ Bad
CREATE TABLE dbo.Employees (
    EmployeeID INT PRIMARY KEY,
    DepartmentID INT,
    DepartmentName NVARCHAR(50) -- violates 3NF
);
```

```
-- ✓ Good
CREATE TABLE dbo.Departments (
    DepartmentID INT PRIMARY KEY,
    DepartmentName NVARCHAR(50) NOT NULL
);

CREATE TABLE dbo.Employees (
    EmployeeID INT PRIMARY KEY,
    DepartmentID INT FOREIGN KEY REFERENCES dbo.Departments(DepartmentID)
);
```

3.6 Practical Tips for Data Engineering

1. Always define PKs — makes joins predictable.
2. Use FKs if you can — catches data issues early.
3. Normalize for OLTP, denormalize for analytics — star schema, fact/dimension tables.
4. Naming conventions matter — e.g., `tbl_Orders` vs `Orders`, `FK_Orders_Customers`.
5. Indexes go hand-in-hand with PK/FK — speed joins and filters.

3.7 Hands-On Exercise

1. Create `Customers`, `Products`, `Orders`, and `OrderItems` tables.
2. Insert sample data:

```
INSERT INTO dbo.Customers (FirstName, LastName, Email) VALUES
('Alice', 'Smith', 'alice@example.com'),
('Bob', 'Jones', 'bob@example.com');

INSERT INTO dbo.Products (ProductName, Price) VALUES
('Widget', 19.99),
('Gadget', 29.99);

INSERT INTO dbo.Orders (CustomerID) VALUES (1), (2);

INSERT INTO dbo.OrderItems (OrderID, ProductID, Quantity, UnitPrice) VALUES
(1, 1, 2, 19.99),
(1, 2, 1, 29.99),
(2, 2, 3, 29.99);
```

3. Query total spending per customer:

```
SELECT c.FirstName, c.LastName, SUM(oi.Quantity * oi.UnitPrice) AS TotalSpent
FROM dbo.Customers c
JOIN dbo.Orders o ON o.CustomerID = c.CustomerID
JOIN dbo.OrderItems oi ON oi.OrderID = o.OrderID
GROUP BY c.FirstName, c.LastName;
```

Expected result: Alice and Bob's total spendings calculated correctly.