

15. Window Functions — Deep Dive

Window functions let you look across rows *without* collapsing them. They're like telescopes you mount on your SELECT clause: you see the neighborhood around each row while still getting one output row per input row.

A window function = `function() OVER (PARTITION BY ... ORDER BY ... <frame clause>)`

Why window functions matter

They solve problems that are painful with subqueries or joins:

- running totals
- percentiles
- rankings and row numbering
- comparing a row to the previous or next row
- computing moving averages
- detecting gaps
- advanced analytics (distribution functions)

They're indispensable in data engineering pipelines when shaping fact tables, deduplicating changefeeds, or building analytic summaries.

The OVER() clause — the control center

`OVER()` defines the *window* of rows a function can see.

- `PARTITION BY` divides rows into groups.
- `ORDER BY` defines the sequence within each partition.
- `ROWS` or `RANGE` defines the frame of rows accessible.

Example structure:

```
SUM(SalesAmount) OVER (
    PARTITION BY CustomerID
    ORDER BY OrderDate
    ROWS BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW
) AS RunningTotal
```

ROW_NUMBER(), RANK(), DENSE_RANK()

ROW_NUMBER()

Assigns a unique sequence per partition.

```
SELECT
    OrderID,
    ROW_NUMBER() OVER (PARTITION BY CustomerID ORDER BY OrderDate) AS Seq
FROM Sales.Orders;
```

Great for deduplication patterns:

```
WITH Ranked AS (
    SELECT *, ROW_NUMBER() OVER (PARTITION BY CustomerID ORDER BY UpdatedAt DESC) AS rn
    FROM Staging.Customers
)
SELECT * FROM Ranked WHERE rn = 1;
```

RANK() and DENSE_RANK()

```
RANK(): 1, 2, 2, 4
DENSE_RANK(): 1, 2, 2, 3
```

LAG() and LEAD() — seeing previous and next rows

Compare a row to its neighbors.

```
SELECT
    OrderDate,
    Revenue,
    LAG(Revenue) OVER (ORDER BY OrderDate) AS PrevRev,
    LEAD(Revenue) OVER (ORDER BY OrderDate) AS NextRev
FROM Finance.DailyRevenue;
```

Useful for detecting shifts, gaps, and change deltas:

```
SELECT
    OrderID,
    Quantity - LAG(Quantity) OVER (ORDER BY OrderID) AS DiffFromPrev
FROM Inventory.Movements;
```

FIRST_VALUE(), LAST_VALUE(), NTH_VALUE()

FIRST_VALUE

```
FIRST_VALUE(Price) OVER (ORDER BY Price ASC)
```

LAST_VALUE

Needs a frame fix; default frames can make it return the current row.

```
LAST_VALUE(Price) OVER (
    ORDER BY Price
    ROWS BETWEEN UNBOUNDED PRECEDING AND UNBOUNDED FOLLOWING
)
```

Aggregate windows — the running total classics

Running totals

```
SUM(Amount) OVER (
    PARTITION BY AccountID
    ORDER BY TxnDate
    ROWS BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW
) AS Balance
```

Moving averages

```
AVG(Temperature) OVER (
    ORDER BY ReadingTime
    ROWS BETWEEN 6 PRECEDING AND CURRENT ROW
) AS SevenPointAvg
```

ROWS vs RANGE — subtle but important

ROWS

Counts physical rows.

```
ROWS BETWEEN 2 PRECEDING AND CURRENT ROW
```

Always deterministic.

RANGE

Groups rows by ORDER BY value equivalence.

```
RANGE BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW
```

Not allowed for non-numeric ORDER BY expressions.

NTILE() — bucketing

Create percentiles or quartiles.

```
SELECT
    CustomerID,
    NTILE(4) OVER (ORDER BY TotalSpent) AS Quartile
FROM Analytics.CustomerSpend;
```

Perfect for segmentation.

Distribution functions: PERCENT_RANK(), CUME_DIST()

PERCENT_RANK

Position of a row in its partition between 0 and 1.

```
PERCENT_RANK() OVER (ORDER BY Score)
```

CUME_DIST

Cumulative distribution; fraction of rows \leq current row.

```
CUME_DIST() OVER (ORDER BY Score)
```

Practical data engineering patterns

1. Deduplicating CDC streams (pick latest version)

```
WITH v AS (
    SELECT *, ROW_NUMBER() OVER (
        PARTITION BY BusinessKey ORDER BY UpdatedAt DESC) AS rn
    FROM Raw.CDC_Events
)
SELECT * FROM v WHERE rn = 1;
```

2. Detecting gaps in sequences

```
SELECT
    EventID,
    EventDate,
    LAG(EventDate) OVER (ORDER BY EventDate) AS Prev,
    DATEDIFF(day, Prev, EventDate) AS GapDays
FROM Logs.Events;
```

3. Building SCD Type 2 boundaries

```
SELECT
    EmployeeID,
    StartDate,
    LEAD(StartDate) OVER (PARTITION BY EmployeeID ORDER BY StartDate) AS EndDate
FROM HR.EmployeeHistory;
```

Produces inferred end dates.

4. Running inventory balances

```
SELECT
    ItemID,
    TxnDate,
    SUM(DeltaQty) OVER (
        PARTITION BY ItemID ORDER BY TxnDate
        ROWS BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW
    ) AS StockLevel
FROM Warehouse.InventoryMoves;
```

Final notes

Window functions don't reduce row count. They augment each row with context. They execute *after WHERE but before ORDER BY* in the logical query order.

This chapter completes the analytic machinery of T-SQL. Next up, if you continue the series, we can explore indexing, transactions, performance tuning, or error handling — wherever your data engineering map leads.